

# Summary

You will design and implement an **Agentic Retrieval-Augmented Generation (RAG)** chatbot that serves as a travel guide by ingesting a travel-focused knowledge base and orchestrating multiple autonomous agents via the **CrewAI** framework. Your system will index and retrieve travel documents (e.g., hotel reviews, itineraries) using a vector database, then generate user-centric recommendations. You'll wrap the chatbot in a **Streamlit** UI and expose it securely over the internet using **ngrok**, enabling live demos. This project demonstrates end-to-end skills in data ingestion, vector search, multi-agent orchestration, and web deployment.

---

## Description

- **Objective:** Build a multi-agent RAG travel guide chatbot that:
    1. Processes and indexes travel documents (e.g., hotel reviews, itineraries).
    2. Retrieves relevant information using vector search.
    3. Coordinates multiple CrewAI agents (retriever, summarizer, response-composer).
    4. Presents a conversational interface via Streamlit and ngrok.
  - **Scope:**
    1. **Data Layer:** Ingest a travel dataset (e.g., Bitext Travel LLM corpus or TripAdvisor hotel reviews).
    2. **RAG Pipeline:** Implement vectorization, indexing (e.g., Qdrant), and retrieval logic.
    3. **Agentic Orchestration:** Configure CrewAI agents for retrieval, filtering, and response crafting.
    4. **Deployment:** Serve the app locally with Streamlit and tunnel it using ngrok.
- 

## Requirements

## 1. Environment Setup

- **Python ≥3.9** with virtual environment.

Install **CrewAI**:

Install **vector database client** (e.g., Qdrant or Pinecone):

Install **Streamlit** and **pyngrok**:

## 2. Data Ingestion

- **Dataset Selection:**
  - Option A: **Bitext Travel LLM Chatbot Training Dataset** (4.16M tokens) on Hugging Face
    - [Hugging Face](#).
  - Option B: **TripAdvisor Hotel Reviews** (20K reviews) on Kaggle[Kaggle](#).
- **Preprocessing:**
  - Clean and normalize text (remove HTML, special characters).
  - Split into document chunks (≈500 tokens each).
  - Serialize chunks with metadata (source, timestamp, category).

## 3. Vector Indexing & Retrieval

- **Vectorization:** Embed each chunk using an open-source embedding model (e.g., OpenAI embeddings or Cohere).
- **Indexing:** Ingest embeddings into Qdrant (or alternative) with semantic metadata filters (e.g., hotel vs. flight data).
- **Retrieval Agent:** Configure a CrewAI agent that accepts a user query and returns the top-k semantically similar chunks

[Qdrant](#).

## 4. Response Generation

- **Summarizer Agent:** Aggregate retrieved chunks, removing redundancy.
- **Composer Agent:** Leverage an LLM (OpenAI GPT-4 or similar) to generate the final user-facing response, formatted as a travel itinerary or recommendation.

## 5. Streamlit Frontend

- Build an interactive UI with:
    - **Text input** for user queries.
    - **Chat window** showing conversation history.
    - **Sidebar** for agent-chain status or logs.
- 

## Deliverables

Component	What to Submit
1. Environment Setup	- <code>requirements.txt</code> listing all dependencies - Setup instructions ( <code>README.md</code> )
2. Data Ingestion	- Jupyter notebook ( <code>data_ingest.ipynb</code> ) showing dataset download, cleaning, chunking - Sample processed JSON or CSV files
3. Indexing & Retrieval	- Script/notebook ( <code>index_and_retrieve.py</code> ) demonstrating embedding and upload to Qdrant - Retrieval demo (unit tests or example queries)
4. Agentic Orchestration	- CrewAI configuration file ( <code>crew_config.yml</code> or Python script) - Agent definitions with comments explaining each role
5. Response Generation	- Python module ( <code>compose_response.py</code> ) invoking the summarizer and composer agents - Example prompts and outputs

- 6. Streamlit App** - `app.py` with Streamlit code<br/>- Screenshot of UI layout<br/>- `ngrok` launch script or instructions
- 7. Documentation & Demo** - **Video walkthrough** (5–7 min) showing end-to-end flow<br/>- **Project report** (PDF) describing design decisions, architecture diagram, and future improvements